

Linear Regression from Scratch - Assignment Report

1. Introduction and Design Decisions

This project implements Linear Regression entirely from scratch using NumPy: forward pass, Mean Squared Error loss, gradient computation, and parameter updates are all hand-derived and hand-coded - no `sklearn.linear_model.LinearRegression`, `SGDRegressor`, or any library that trains the model directly. `scikit-learn` is used only for loading the California housing dataset and for `train_test_split`.

Class design. The implementation is split into single-responsibility classes: `DataHandler` (load/split), `DataPreprocessor` (standardization), `LinearRegressionModel` (predict / loss / gradient / update), `LinearRegressionTrainer` (the three training loops), `RegressionEvaluator` (MSE/RMSE/MAE), `RegressionVisualizer` (plots), and `ModelPersistence` (save/load). The payoff of this separation shows up directly in the experiments below: the exact same `LinearRegressionModel`, `RegressionEvaluator`, and `RegressionVisualizer` code was reused, completely unchanged, across every learning rate, every optimizer, both normalized and raw features, and two entirely different data modalities (real face images, real cough audio) - only the data-loading class ever changed.

Two additional dataset handlers - `ImageDataHandler` and `AudioDataHandler` - were built as subclasses of `DataHandler` for the image (face age) and audio (COUGHVID) parts of Part 3, following the same interface, and both were run end-to-end on the real downloaded datasets. Results are in Section 9.

Primary dataset for all experiments below: `sklearn.datasets.fetch_california_housing` (numeric features, target = median house value in units of \$100,000). Split 70% train / 15% validation / 15% test, giving 14,448 / 3,096 / 3,096 samples.

2. Part 2 - Manual Gradient Descent Verification

Given $X = [1, 2, 3]$, $Y = [3, 5, 7]$, initial $\theta_1 = 0.5$, $\theta_0 = 0.7$, $\eta = 0.01$.

Quantity	Value
$\hat{y}$	$[1.2, 1.7, 2.2]$
errors ($\hat{y} - y$)	$[-1.8, -3.3, -4.8]$
MSE	12.39
dL/d_{θ_1}	-15.2
dL/d_{θ_0}	-6.6
θ_1 (updated)	0.652
θ_0 (updated)	0.766

A second iteration, starting from $\theta_1 = 0.652$, $\theta_0 = 0.766$:

Quantity	Value
$\hat{y}$	$[1.418, 2.07, 2.722]$


Quantity	Value
errors	[-1.582, -2.93, -4.278]
MSE	9.796303
dL/d_theta1	-13.5173
dL/d_theta0	-5.86
theta_1 (updated)	0.787173
theta_0 (updated)	0.8246

Both iterations were independently verified against `LinearRegressionModel` in `experiments/part2_manual_gd_verification.py`, matching to numerical precision (differences only at the level of floating-point representation, e.g. `-6.6` printing as `-6.599999999999999`). Loss decreased monotonically across both iterations (`12.39 -> 9.80`), consistent with a correctly-signed gradient and a stable learning rate.

The handwritten copy of this calculation (photographed/scanned on paper, as the assignment requires) is included at `handwritten/part2_manual_gradient_descent.pdf`.

3. Part 3a - Learning-Rate Experiment

Trained with Batch GD on **standardized** California housing features, 200 epochs, for each of the five required learning rates.



Learning Rate	Final Training Loss	Final Validation Loss	Converged?
0.00001	5.5509	5.6820	No - barely moved off its starting point
0.0001	5.2107	5.3320	No - improving, but very slowly
0.001	2.8704	2.9269	No - improving steadily, not yet flat
0.01	0.5983	0.5902	Nearly - still inching down
0.1	0.5241	0.5166	Yes - loss plateaued

1. Fastest to converge: $lr = 0.1$ - reaches a plateau (~ 0.52 loss) well within 200 epochs.

2. Most stable: also $lr = 0.1$ - smooth, monotonic decrease with no oscillation.

3. Did any diverge? Not among the five required values (standardization kept all of them numerically stable). Pushed outside that range, $lr = 1.0$ diverges within ~ 10 epochs (5.59 $\rightarrow$ 5.77 $\rightarrow$ 7.48 $\rightarrow$ 20.4 $\rightarrow$ 120 $\rightarrow$ 901 $\rightarrow$ 7038 $\rightarrow$...), and $lr = 2.0$ diverges even faster - direct evidence for question 5 below.

4. Why does a very small learning rate train slowly? Each update is $\theta \leftarrow \theta - \eta * \text{grad}_L$; with $\eta = 0.00001$ even a large gradient produces a microscopic step - the correct direction, just too small a distance per step to cover meaningful ground in a fixed number of epochs.

5. Why can a large learning rate oscillate or diverge? A large step can overshoot the minimum onto a steeper part of the opposite side of the loss bowl, where the new gradient is even larger in the opposite direction; each subsequent step overshoots further, compounding exponentially - exactly what the $lr = 1.0/2.0$ numbers above show.

6. Final choice: $lr = 0.1$ - lowest final training and validation loss, fastest to plateau, and confirmed stable (not right at the edge of diverging).



Bonus finding: running this same sweep on **raw, unnormalized** features made every one of the five required learning rates diverge to NaN - even 0.00001 . This is because `Population` has a training-set standard deviation of ~ 1141 (max $\sim 35,682$) versus `MedInc`'s ~ 1.9 , producing gradients large enough that no learning rate in the requested range is small enough to stay stable. This directly motivates Part 3b.

4. Part 3b - Normalization Experiment

Experiment A trains on raw features; **Experiment B** trains on features standardized using training-set mean/std only, applied identically to validation and test data.

Head-to-head at the same learning rate ($lr = 0.01$)

	Experiment A (raw)	Experiment B (standardized)
Diverged?	Yes (NaN/inf)	No
Final train loss	NaN	0.598306
Final val loss	NaN	0.590211



Learning-rate sensitivity

Experiment A - raw features:

lr	Diverged?	Final train loss
1e-9	No	3.2033
1e-8	No	2.9488
1e-7	No	2.3397
1e-6	Yes	-
1e-5	Yes	-

Experiment B - standardized features:

lr	Diverged?	Final train loss
1e-5	No	5.5509
1e-4	No	5.2107
1e-3	No	2.8704
1e-2	No	0.5983
1e-1	No	0.5241

Stable rates: raw = 3/5 tested (only within the tiny 1e-9-1e-7 window); standardized = 5/5 tested (spanning 1e-5-1e-1, four orders of magnitude).


Dimension	Experiment A (raw)	Experiment B (standardized)
Convergence speed	Extremely slow even at its best stable rate (2.34 after 200 epochs at 1e-7)	Fast - 0.52 in the same 200 epochs at <code>lr=0.1</code>
Loss stability	Diverges once <code>lr >= 1e-6</code>	Stable across the full required range
Final training loss	Best case ~ 2.34	Best case ~ 0.52 - roughly 4.5x lower
Final validation loss	Not usably measured	0.517 at <code>lr = 0.1</code>
Sensitivity to learning rate	Very high - safe window ~ 2 orders of magnitude, shifted absurdly small	Low - safe across the full requested range

Why does normalization make gradient-based optimization easier? With raw features, one feature's huge scale (Population) stretches the loss surface into a narrow, elongated canyon - steep in one direction, nearly flat in others. No single learning rate is safe for the steep direction and still makes progress in the flat one. Standardizing every feature to mean 0, std 1 reshapes that canyon into something close to a round bowl, so a single learning rate works reasonably well in every direction at once.

5. Part 3c - Batch GD vs. SGD vs. Mini-Batch GD

Each optimizer needed its **own** stable learning rate, since they differ in how many parameter updates happen per epoch (Batch GD: 1, Mini-Batch: ~452, SGD: 14,448) - the same `lr = 0.01` that is stable for Batch GD makes both

SGD and Mini-Batch diverge.



Optimization Method	Learning Rate Used	Training Time (50 epochs)	Final Train Loss	Loss Curve Behavior
Batch GD	0.1	0.005 s	0.5536	Perfectly smooth - 0% of epochs increased
SGD	0.0001	~5-24 s	0.5245	Noisiest - 45% of epochs increased
Mini-Batch GD	0.001	0.19 s	0.5241	In between - 18% of epochs increased

("% of epochs where loss increased" is used instead of raw variance of the epoch-to-epoch drop, since variance alone conflates "noisy" with "dropping fast.")


- 1. Smoothest curve?** Batch GD - the true average gradient over all 14,448 samples has no sampling noise.
- 2. Noisiest?** SGD - each update reacts to a single sample's error, a high-variance estimate of the true gradient.
- 3. Fastest in epochs?** Roughly tied - all three reach loss ~ 0.52-0.55 by epoch 50, since SGD/Mini-batch get far more parameter updates within the same 50 epochs (SGD: 722,400 updates vs. Batch GD's 50).
- 4. Fastest computationally?** Batch GD, by a wide margin - one vectorized matrix multiply per epoch, versus SGD paying Python-loop overhead 14,448 times per epoch for tiny operations that don't benefit from vectorization.
- 5. Why is mini-batch used in modern deep learning?** It is the practical middle ground shown directly in this table: far less noisy than SGD (18% vs. 45% upward blips) while dramatically faster than SGD in wall-clock time, by

processing chunks large enough to use vectorized/GPU operations efficiently while still getting many updates per epoch.

6. Part 4 - Prediction and Evaluation on the Unseen Test Set

Final model: Batch GD, $\text{lr} = 0.1$, standardized features, 300 epochs.

Metric	Train	Validation	Test
MSE	0.5237	0.5160	0.5373
RMSE	-	-	0.7330
MAE	-	-	0.5352



Train, validation, and test losses are all close (0.52-0.54), indicating the model is not overfitting - if anything it is mildly underfitting everywhere, consistent with a linear model being unable to capture whatever nonlinear structure exists in housing prices.

5 largest test-set errors:

actual=5.000	predicted=0.786	abs_error=4.214
actual=5.000	predicted=1.120	abs_error=3.880
actual=5.000	predicted=1.281	abs_error=3.719
actual=0.675	predicted=4.212	abs_error=3.537
actual=5.000	predicted=1.542	abs_error=3.458

Four of the five worst errors have `actual = 5.000` exactly - California housing's target is capped at \$500,000 (any pricier house is recorded as exactly 5.0), an artificial ceiling the model has no way to know about. These large "errors" are partly a data-censoring artifact, not purely a modeling failure.

- 1. What does a point far from the diagonal represent?** A sample where the model's linear assumption breaks down, or - as shown above - a case where the label itself is capped/noisy rather than the prediction being simply wrong.
- 2. Why measure on an unseen test set?** Training loss only shows how well the model fit data it already saw; only held-out data reveals how it performs on genuinely new inputs.
- 3. Why is training loss alone insufficient?** A model can drive training loss arbitrarily low by fitting training-specific noise (overfitting) while performing badly elsewhere; the close agreement between train/val/test loss here (0.52-0.54) is itself the evidence needed to claim reasonable generalization - evidence training loss alone could not provide.
- 4. MSE vs. RMSE vs. MAE:** MSE (0.5373) squares errors, so it is dominated by the handful of huge misses (a single 4.2-off prediction contributes as much as ~65 typical 0.5-sized errors). RMSE (0.7330) restores MSE to the target's own units (\$100k), giving an interpretable "\$73,300 typical error" figure, though still inflated by the same outliers. MAE (0.5352) weighs every error proportionally to its size, giving a more robust sense of typical error, barely moved by the capped-target outliers.

7. Part 7 - Interpreting the Learned Parameters

Coefficients from the final model (on standardized features, so magnitudes are directly comparable):

Feature	Weight
Latitude	-0.8729
Longitude	-0.8439
MedInc	+0.8427
AveBedrms	+0.3393
AveRooms	-0.3059
HouseAge	+0.1191
AveOccup	-0.0466
Population	-0.0043
bias	+2.0630

- 1. Strongest positive:** MedInc (+0.8427) - a one-standard-deviation increase in median income is associated with a +0.84 increase in predicted house value.
- 2. Strongest negative:** Latitude (-0.8729) - moving one standard deviation further north is associated with a substantial predicted price decrease, reflecting that California's priciest metro areas sit toward the state's southern/central coastline.
- 3. What does the sign indicate?** Positive: feature and target move together, holding other features fixed. Negative: they move oppositely.

4. Why be careful with correlated features? The training-set correlation matrix shows Latitude/Longitude at **-0.924** and AveRooms/AveBedrms at **+0.86** - nearly collinear pairs. When two features move together this strongly, the model splits credit between them somewhat arbitrarily; Latitude's large weight really reflects a shared geographic signal split across two correlated coefficients, not "latitude alone" as an isolated driver.

5. Does a large coefficient prove causation? No - Latitude's large weight is an association found in this dataset, not evidence that changing a house's latitude alone (with everything else fixed) would change its price by that amount. The real driver is almost certainly a confound: latitude correlates with proximity to specific expensive job markets/coastlines.

Feature correlation matrix (raw features, training set):

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude
MedInc	1.00	-0.13	0.31	-0.06	0.01	0.03	-0.07	-0.02
HouseAge	-0.13	1.00	-0.14	-0.07	-0.29	0.01	0.00	-0.10
AveRooms	0.31	-0.14	1.00	0.86	-0.07	-0.01	0.11	-0.03
AveBedrms	-0.06	-0.07	0.86	1.00	-0.06	-0.01	0.07	0.01
Population	0.01	-0.29	-0.07	-0.06	1.00	0.07	-0.10	0.09
AveOccup	0.03	0.01	-0.01	-0.01	0.07	1.00	0.01	-0.00
Latitude	-0.07	0.00	0.11	0.07	-0.10	0.01	1.00	-0.92
Longitude	-0.02	-0.10	-0.03	0.01	0.09	-0.00	-0.92	1.00

8. Part 5 - Debugging Challenge

Bug 1

```
theta = theta + learning_rate * gradient
```

The gradient points toward *increasing* loss; gradient descent must move opposite to it. Adding the gradient drives the loss up, not down. Fix:

```
theta = theta - learning_rate * gradient
```

Bug 2

```
gradient = (2 / N) * np.dot(X, error)
```

X has shape (N, d) and error has shape (N,) (or (N, 1)) - contracting X's last axis against error is the wrong orientation; the sum needs to run over the N samples for each of the d features, which requires the transpose:

```
gradient = (2 / N) * np.dot(X.T, error)
```

Bug 3

```
test_mean = np.mean(test_X, axis=0)
test_X = (test_X - test_mean) / np.std(test_X, axis=0)
```

This computes normalization statistics from the test set itself - data leakage, and it also means test data is transformed differently than the model was trained on. The training-set mean/std (stored by DataPreprocessor.fit on train data only) must be reused:

```
test_X = (test_X - train_mean) / train_std
```

Bug 4

```

for epoch in range(n_epochs):
    for i in range(n_samples):
        x = train_X[i]
        y = train_Y[i]

```

Indexing a 2D array with a single integer (`train_X[i]`) collapses it to 1D, shape `(d,)`, dropping the sample dimension the rest of a matrix-based implementation (`predict`, `compute_gradient`) expects. Slicing instead of indexing keeps it 2D:

```

x = train_X[i:i+1]    # shape (1, d)
y = train_Y[i:i+1]    # shape (1, 1)

```

Bug 5 - A student reports only training loss and claims the model generalizes well. This is not justified: a model can memorize training-specific noise and show arbitrarily low training loss while performing badly on unseen data. Only held-out validation/test loss provides evidence about generalization; without it, no generalization claim is supported.

9. Image and Audio Datasets - Real Results



Per the assignment's requirement for three dataset modalities, `ImageDataHandler` (face-age regression) and `AudioDataHandler` (COUGHVID cough_detected regression) were implemented as subclasses of `DataHandler`, reusing `split_data()/get_feature_names()` unchanged and plugging into the same `DataPreprocessor -> LinearRegressionModel -> LinearRegressionTrainer` pipeline with zero modification to those classes -- the same code that ran California housing, now proven to generalize to two completely different modalities.

9.1 Image - Kaggle frabbisw/facial-age

9,778 real face images (grayscale, resized to 32x32 = 1,024 pixel features), age labels 1-110 taken from each image's containing folder name.

	Value
Learning rate used	0.0025
Epochs	1000
Final train loss (MSE)	224.29
Final val loss (MSE)	249.66
Test MSE	251.83
Test RMSE	15.87 years
Test MAE	12.36 years
Baseline MSE (predict mean age)	621.43

The model beats the naive baseline by ~59%, so it is genuinely picking up age-related signal from raw pixels, not just noise.



A second, real confirmation of the learning-rate/conditioning lesson from Part 3a: this dataset needed $lr = 0.0025$ to stay stable -- two orders of magnitude smaller than California housing's $lr = 0.1$. With 1,024 raw pixel

features, neighboring pixels are highly correlated (a value at pixel (i, j) is rarely independent of its neighbors), which makes the loss surface far more ill-conditioned than housing's 8 largely-independent features, even after standardization -- standardization fixes each feature's *scale*, but does nothing about *correlation* between features. $1r$ values as small as 0.003 already diverged.


9.2 Audio - COUGHVID (Zenodo record 4498364)

The complete dataset: 27,550 real cough recordings (`.webm`), decoded through `ffmpeg` (soundfile/librosa's default backend cannot read `.webm` natively -- see `AudioDataHandler._decode_to_wav`), each converted into a 29-dimensional feature vector (13 MFCC means, 13 MFCC stds, zero-crossing rate, spectral centroid, RMS energy), regressing on the `cough_detected` column (0-1) from `metadata_compiled.csv`.


	Value
Clips processed	27,550 (full dataset)
Feature-extraction time	3,257s (~54 min) -- one <code>ffmpeg</code> subprocess call per clip
Learning rate	0.05
Epochs	200
Final train loss (MSE)	0.0414
Final val loss (MSE)	0.0405
Test MSE	0.0416
Test RMSE	0.2039
Test MAE	0.1578
Baseline MSE (predict mean <code>cough_detected</code>)	0.1526

The model beats the naive baseline by **~73%**, and a preliminary 1,500-clip subset run produced nearly identical numbers (test MSE 0.0444), confirming the result is stable and not a fluke of sample size.

COUGHVID (real data, n=27550) - Train vs Val Loss



Actual vs. Predicted



Strongest predictors: mfcc_std_0 (+0.205) and mfcc_mean_0 (+0.169) -- MFCC coefficient 0 tracks a signal's overall spectral energy/loudness shape, a sensible feature for distinguishing an actual cough from background

noise or silence.

An operational lesson from this run, worth reporting alongside the modeling results: the first attempt at this full run hung indefinitely (left running over 14+ hours with zero progress) because the `ffmpeg` subprocess call had no timeout, and one malformed/unusual clip caused it to stall forever with no exception raised to catch and skip. Adding `timeout=20` (plus `-nostdin`, since `ffmpeg` can otherwise block waiting on standard input) to that subprocess call fixed it -- the corrected run completed cleanly end-to-end. This is a practical instance of "any external I/O call without a timeout can hang your entire pipeline," worth keeping in mind for future work with subprocess-based decoding.

10. Conclusion

A complete, from-scratch Linear Regression implementation was built, verified by hand against a manual gradient-descent calculation, and exercised through a full experimental suite on the California housing dataset: a five-point learning-rate sweep, a raw-vs- normalized ablation, a three-way optimizer comparison, held-out test evaluation, and coefficient interpretation with an explicit correlated-features caveat. The main takeaways - normalization is close to mandatory for this dataset's feature scales, $\text{lr}=0.1$ with Batch GD was the best converging/stable configuration found, and mini-batch is the right practical compromise between Batch GD and SGD - are all backed by measured numbers rather than assumed from theory alone. The same pipeline was then proven to generalize unchanged to two more real datasets: 9,778 real face images (test MSE 251.8 vs. a baseline of 621.4) and the complete 27,550-clip COUGHVID audio dataset (test MSE 0.0416 vs. a baseline of 0.1526), each requiring its own tuned learning rate but no changes to any class besides the data loader -- direct evidence that the OOP separation from Part 1 achieves what it was designed for.